

Report: Project - Part 2

Martijn Leplae (*r0737706*)

Andreas Hinderyckx (*r0760777*)

December 17, 2021

Question 11 NoSQL affects both elements as follows:

1. Scalability:

- NoSQL is easily scalable: instead of needing to upgrade the hardware (i.e. CPU) one can simply add extra servers to distribute the load.
- New data can be incorporated into the database more easily, instead of needing to redesign the entire schema of an SQL-database.
-

2. Availability:

- The data is then duplicated, which increases redundancy to mitigate the effect of crashes or failures of certain servers.
- It can be spread geographically, thus making it easy to make available to lots of places in the world.
- Data can be stored 'as is', that is: it needs not to be transformed in order to be put into the database, but can rather be inserted using its original data types.

Question 12 In our Firestore database we have set up three collections: "LocalShows", "Bookings", "BestCustomers". As primary keys, UUIDs of the shows and bookings are used for the first two collections, and the customer's email address is used in the customers collection.

- In **BestCustomers** we keep track of the number of tickets per customer. This facilitates efficient querying for the **getBestCustomer** method, as we do not need to scan every customer's list of bookings.
- The collection **Bookings** is used to store the bookings.
- **LocalShows** holds the local shows. However, we opted to set up two subcollections here: **Seats** and **showTimes**. We do this for the following reasons:
 1. It makes the querying of an individual Seat way easier and faster E.g when querying a specific seat of a show, no other data needs to be send over the net, all the querying happens in Firestore. This prevents the combinatorial explosion of looping over all possible seat-time combinations.
 2. The showtimes could be derived from the seats. However, in that case we would need to download all the seats and do a lot of querying ourselves. For this reason, having a separate collection for the Showtimes is worth the small amount of data duplication in return for efficient querying.

Question 13 NoSQL offers a much more limited querying interface. It is not possible to link data from different collections together using the NoSQL querying interface. To resolve this, a new collection must be made containing these queried elements. In the case of SQL, one can use more intricate queries without needing to make modifications to the data model.

Question 14 In the case of all-or-nothing semantics with the remote companies, the DELETE-request made to undo a failed booking might fail too, whereas using a transaction guarantees that its effects can be rolled back when a failure occurs.

With our previous implementation, we sequentially booked tickets one-by-one and rolled back one-by-one on failure. The transaction on the other hand first makes sure that all the seats can be booked (i.e. they are not already booked), and then makes all the bookings.

Question 15 The e-mail feedback is necessary as Pub-Sub handles the booking confirmation process, which means that this might incur some delay and that the web-interface won't be able to instantly show the confirmation of the booking to the user. Therefore, when a booking is made a confirmation e-mail is sent to the client to either confirm his booking or to report that his booking didn't succeed. Based upon this information the client may re-submit his booking in case of failure or may be reassured of the success of his booking in case of success.

Question 16 In case of peak usage of the website — e.g. when there would be a lot of interest in a certain show at a specific time — we could prepare for this by reducing the `max_pending_latency` and `min_pending_latency` to make sure that enough instances will be available to accommodate for the peak usage. We only changed the `min_instances` to a small positive value to minimise the loading times for the first visitor (as there is then no extra waiting time incurred for starting up the first instances). The other default values seemed to be fine since we don't expect lots of traffic. In essence there will always be a trade-off between making the scaling more responsive on one hand, and limiting costs on the other hand.

Question 17

1. The advantages of running an application as a web service are:
 - Reduced hardware needs: the application can be ran from any device with an internet connection, rather than being limited to devices with hardware sufficiently powerful to run the application,
 - Increased fault-tolerance: local applications only have a single copy of the stored information, whereas a web service can leverage on the advantages of a distributed setting where the information is duplicated across multiple geographically different locations. This leads to overall increased availability,
 - Independence of the hardware the application is run on: native applications need to be ported to every supported hardware and operating system configuration, whereas an application running as a server is independent of these factors,
 - The software is automatically up-to-date, which means the users don't need to do updates themselves. (e.g. there are still people using old software that doesn't have recent security patches for example).
2. Deploying the project in the cloud has the advantages of leveraging on Google's large server infrastructure, which allows for seamless scaling in combination with the automatic allocation of extra instances when needed. This also eliminates the hassle that comes with your own server infrastructure. Besides, the cost of setting up a remote-accessible server is much lower with Google's pay-as-you-go cost model, in comparison to the cost and time involved in acquiring your own hardware and setting this up.

Question 18 Pitfalls of migrating a local project to the cloud include:

- Correctly implementing security,
- Learning to work with and link together the broad spectrum of APIs used in a cloud environment,
- You become dependent of the ecosystem your cloud provider offers: it may not be easy to change this without modifying a lot of code and/or configuration details, as well as missing out on certain features that are offered by one provider but aren't by another provider.

Restrictions of Google Cloud include;

- Google Cloud allows for third party token verification when using Firestore. This increases flexibility and customisation, at the risk of improper security implementations,
- Google Cloud lots of different APIs which must be correctly tied together in a project,
- Debugging becomes more time consuming, as the logs are only visible after a noticeable delay, and the source code itself can no longer be stepped through to debug potential issues. Deploying takes around a minute, so when you want to see the effect of a (potentially small) change on the cloud deployment, you need to wait for a while.

Question 19 A lot of the code we wrote is specific to make use of the Google Cloud APIs, which means that this all of this code must be rewritten when migrating to a different cloud provider. For example, storing data in the Firestore database required code to specifically connect to this database, create specific data structures (**collections** and **documents**), retrieve this data in a specified manner, etc. This also holds for authentication, pub-sub, etc. So we would need to find alternatives for these on the other provider. Additionally, certain providers may offer a different set of features or APIs which won't be offered by different providers, potentially requiring you to re-design or re-develop certain functionalities.

Question 20 We simulated multiple concurrent visitors of the website by opening 8 different browser windows and using the page (approximately) concurrently. Assuming that the users visit the website uniformly across the 24 hour time-span of one day, this results in a total of two instances being spawned which gives us an estimate of two users that can be handled per instance. We simulated a use case of an average user who books several seats at different shows and recorded the number of read and write operations made in Firestore. This turns out to be approximately 280 read, 5 write and 0 delete operations. Extrapolating this to 1000 daily users, we get roughly 280000 read, 5000 write and 0 delete operations (as the user has no possibility of deleting his bookings). The current total amount of data used is 0.02GB, which is mainly dominated by the storage to store the local shows. The bookings themselves only require a limited amount of storage, thus we estimate that scaling this up to 1000 daily users will amount up to 0.5GB. Debugging our Pub Sub code shows that this same use case creates a message of size approximately 650 bytes. This gives us a total estimate of 0.065 MB of Pub Sub messages per issued per day. Inputting all this information in the Google Cloud Pricing Calculator (<https://cloud.google.com/products/calculator/#id=1781479a-9372-4643-922f-09bbfefa4b2e>), gives us a total monthly cost estimate of 871.43 dollars per month.

Increasing the daily number of users per day to one million, we analogously retrieve a monthly cost estimate of 923,345 dollars per month (based on our rough estimates). As the major part of this monthly cost comes from Google App Engine, this cost can be reduced by reducing the minimum and maximum number of instances, increasing pending latencies and reducing the maximum number of concurrent requests in the `app.yaml` configuration file. At this scale of number of daily users, however, it would be more cost effective to use regular server APIs of Google rather than App Engine which don't incur costs for automatic scaling, etc. Possibly, setting up an own server infrastructure might be more cost effective in the long run, too, at this scale.